

Topic 8: Applied Math (C)

Optimization Methods

Group Number 9

Ryan Hammonds, Benjamin Pham, Gabriel Riegner
08 November 2022

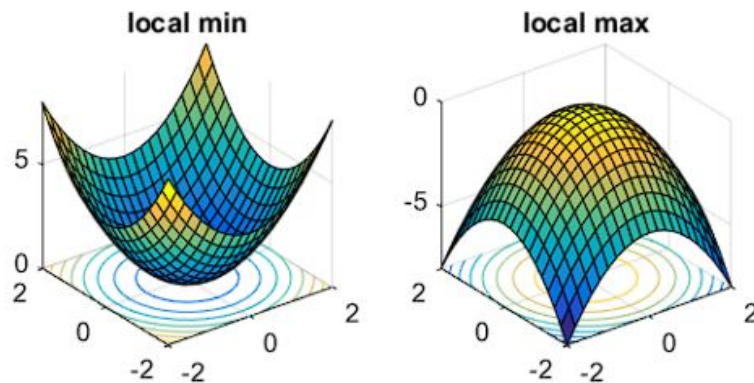
Sec 1: Introduction

Overview of Optimization

- Optimization involves minimizing (or maximizing) an objective or loss function.
- Includes optimizing unknowns in regression and classification problems to minimize error.

$$\min_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x})$$

- Solved using:
 - Linear algebra:
 - $\mathbf{Ax} = \mathbf{b}$
 - Iterative methods
 - Gradient Descent
 - Newton's Method



$$\nabla f(\mathbf{x}) = 0$$

Importance of the topic


- Optimization is required to solve regression and classification problems.
 - Parameter tuning
- Applies to simple and complex models.
 - Linear Regression
 - Neural networks
- Applications
 - Medicine
 - Economics / Finance
 - Computer vision
 - Speech recognition

$$Y_1 = \beta_0 + \beta_1 X_1 + \epsilon_1$$

$$Y_2 = \beta_0 + \beta_1 X_2 + \epsilon_2$$

$$\vdots \quad \vdots \quad \vdots$$

$$Y_n = \beta_0 + \beta_1 X_n + \epsilon_n$$



Sec 2: Problem Formulation

#1 Problem formulation

1. Let \mathbf{x} be a vector of unknown parameters.
2. Let \mathbf{y} be a vector of the known targets.
3. Given an arbitrary loss function (e.g. L0, L1, or L2), iteratively minimize:

$$f(\mathbf{x}) = \text{loss}(\hat{\mathbf{y}}, \mathbf{y})$$

$$\min_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x})$$

4. At each iterative step k , update \mathbf{x} by subtracting either:
 - a. the gradient ∇ scaled by step size η

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \eta_k \nabla f(\mathbf{x}_k)$$

- b. the gradient ∇ scaled by the inverse Hessian \mathbf{H}^{-1}

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \mathbf{H}(\mathbf{x}_k)^{-1} \nabla f(\mathbf{x}_k)$$

#2 Relation to Numerical Linear Algebra

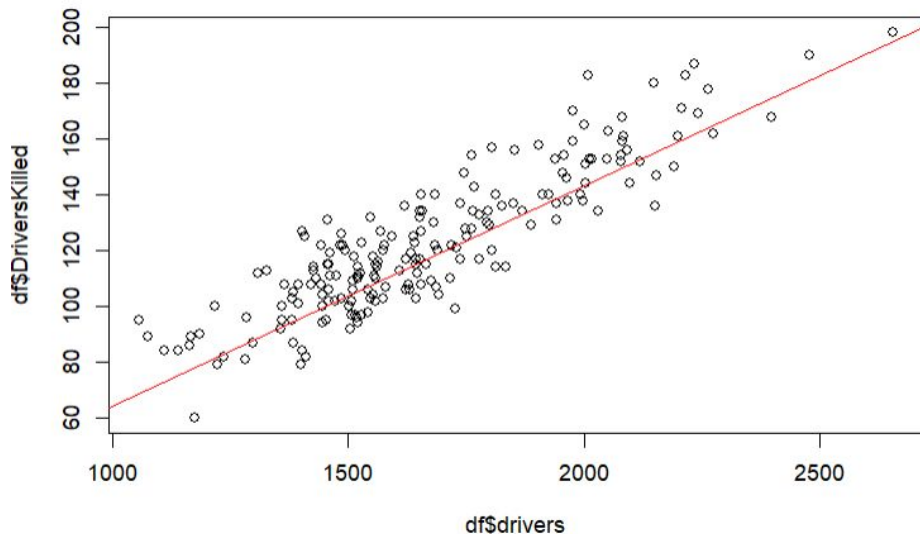
In general, an optimization problem involves **solving a linear system** with multiple parameters to minimize the loss function.

Solved via two main types of methods:

- Closed Form (Exact)
- Numerical (Estimation)

The solution of this linear system results in the set of the most optimal parameters

Simple Example





$$Y = \beta X + \epsilon$$

$$\underbrace{\begin{bmatrix} Y_1 \\ Y_2 \\ Y_3 \\ \vdots \\ Y_n \end{bmatrix} = \begin{bmatrix} 1 & x_{11} & x_{12} \\ 1 & x_{21} & x_{22} \\ 1 & x_{31} & x_{32} \\ \vdots & \vdots & \vdots \\ 1 & x_{n1} & x_{n2} \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \end{bmatrix}}_{\text{Known}} + \underbrace{\begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \epsilon_3 \\ \vdots \\ \epsilon_n \end{bmatrix}}_{\text{Unknown}}$$

Optimization Problem: Get β that creates the **best fit line** while minimizing the error term (distance between points and the best fit line)

Slightly More Complex Example

9



$$\begin{bmatrix} 1 & x_1 & x_2 \end{bmatrix} \begin{bmatrix} 6 & 0 & 0 \\ -2 & 1 & 0 \\ -2 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 6 - 2x_1 - 2x_2 & x_1 & x_2 \end{bmatrix} \begin{bmatrix} 1 \\ x_1 \\ x_2 \end{bmatrix} = 6 - 2x_1 - 2x_2 + x_1^2 + x_2^2 = f(x_1, x_2)$$

Optimization helps find the pair that minimizes this function

#3 Approach of Numerical Linear Algebra (NLA)


- Sometimes, the closed form solution is not available/too difficult to compute

For example: in linear regression, solving for β analytically can be **computationally expensive** in high dimensional design matrices

- Iterative methods such as **Stochastic Gradient Descent** and **Newton's Method** are used instead.

$$\hat{\beta} = \boxed{(X^T X)^{-1}} X^T y$$

Difficult to compute at high dimensions!

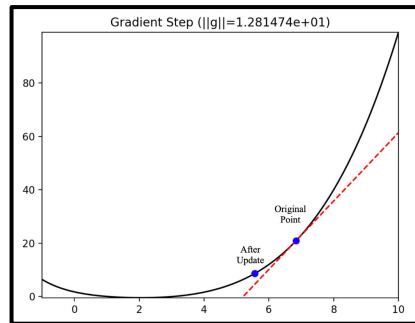



Iterative Optimization Algorithms

$$H = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_d} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_d} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_d \partial x_1} & \frac{\partial^2 f}{\partial x_d \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_d^2} \end{bmatrix}$$

Newton's Method

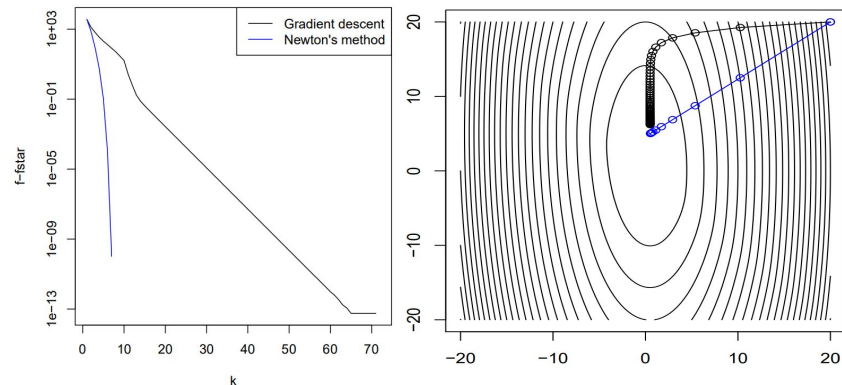
- second order optimization
- fast and precise convergence (error decreases quadratically)
- fewer steps, but each step requires n^2 to compute Hessian and n^3 to invert it
- Hessian matrix might not be invertible (positive definite)



(Stochastic) Gradient Descent

- first order optimization
- direction of most rapid decrease, but no information about curvature
- slow and imprecise (error decreases linearly)

comparison: k steps to convergence



Sec 3: State of the Art (SOTA)

SOTA: Quasi Newton Methods

- Computing the inverse Hessian is costly
- Replace Hessian with an approximation using Taylor expansion

$$\mathbf{B}_k^{-1} \approx \mathbf{H}(\mathbf{x}_k)^{-1}$$

- Hessian approximations must satisfy:

$$\mathbf{B}_{k+1} \Delta \mathbf{x} = \nabla f(\mathbf{x}_{k+1}) - \nabla f(\mathbf{x}_k) \quad \Rightarrow \quad \mathbf{A} \mathbf{x} = \mathbf{b}$$

- Methods to solve \mathbf{B}_{k+1}
 - BFGS
 - Broyden
 - SR1

Sec 4: Experiment

Experimental setup

libraries and tools:





```
def newtons_method(g, max_its, w, **kwargs):
    ... gradient = gradient(g)
    ... hessian = hessian(g)
    ...
```

final report:



datasets/benchmarks:

simulations



sklearn.datasets



	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6				
4	5.0				

	subject	timepoint	event	region	signal
0					
1	s13	18	stim	parietal	-0.017552
2	s5				
3	s12				
4	s11				
5	s10				


	species	island	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g
0	Adelie	Torgersen	39.1	18.7	181.0	3750.0
1	Adelie	Torgersen	39.5	17.4	186.0	3800.0
2	Adelie	Torgersen	40.3	18.0	195.0	3250.0
3	Adelie	Torgersen	NaN	NaN	NaN	NaN
4	Adelie	Torgersen	36.7	19.3	193.0	3450.0

Sec 5: Concluding remarks

Conclusion

Next steps

- **simulation-based optimization** to benchmark linear algebra, gradient descent, and (quasi)-Newton methods



- **evaluation metrics**: **total step number**, **runtime**, **accuracy** with an increasing number of samples and parameters
- Aim to compare performance of each optimization algorithm as number of parameters increase from 1 to 1000
- apply same benchmarks under real data conditions, including regression and classification problems
- **Smallest dataset**: Linnerud (**20 observations, 3 features**)
- **Largest dataset**: Breast Cancer Wisconsin (Diagnostic) dataset (**569 observations, 30 features**)

Expected outcomes

- As the number of features/parameters increases, the step-to-convergence for all algorithms will increase resulting in a longer runtime
- gradient descent will be efficient for small problems, and state-of-the-art methods will be more efficient for high-dimensional, large data examples
- The size of the problem will dictate the optimization algorithm

References

<https://course.ece.cmu.edu/~ece739/lectures/18739-2020-spring-lecture-08-second-order.pdf>

<https://www.psychologie.uni-heidelberg.de/ae/meth/team/mertens/blog/hessian.nb.html>

<https://www.inf.ed.ac.uk/teaching/courses/irds/miniproject-datasets.html>

Hardt, M., & Recht, B. (2022). Patterns, predictions, and actions: A story about machine learning. Princeton University Press.

Boyd, S., & Vandenberghe, L. (2004). Convex optimization. Cambridge university press.